

Penguins

February 27, 2023

0.1 Problem/Question

According to the Australian Antarctic Program (2022), there are a total of 18 species of penguins currently known, 5 of which live in Antarctica, and 4 others live on sub-Antarctic islands. For researchers, it is very important to differentiate among species to better understand mating patterns, eating cycles, and the overall development of penguins. Therefore, we believe it is important to leverage machine learning techniques to distinguish among species, and effectively predict the species of a penguin given certain characteristics.

0.2 Data Sources/References

The data used in this project was collected, cleaned, and published by Dr. Kristen Gorman and the Palmer Station, Antarctica LTER, a member of the Long Term Ecological Research Network. The dataset counts with 7 variables, and a total of 344 observations. All the variables are explained in detail later in the report.

Three main species of penguins are analyzed. Adélie penguins, with a population of 2.4 million, and a lifespan of 15-20 years. Gentoo penguins, with a population of 400,000, and a lifespan of 12-15 years. Lastly, Chinstrap penguins, with a population of 8 million, and a lifespan of 12-15 years.

References:

Australian Antarctic Program (2022) About Antarctica, Animals, Penguins.

<https://www.antarctica.gov.au/>

Gorman KB, Williams TD, Fraser WR (2014) Ecological Sexual Dimorphism and Environmental Variability within a Community of Antarctic Penguins (Genus *Pygoscelis*). PLoS ONE 9(3): e90081. doi:10.1371/journal.pone.0090081

The dataset has 7 columns describing the three species of Penguins. The goal of the project is to create Machine Learning models to predict the Species of the Penguins, and in a similar way to predict the Sex as well. We would like to compare the performance of different models built for predicting each target variable. The models would also be helpful in determining the influential features if any, in determining the Species and Sex of the Penguins. It would be interesting to see if there are any similarities

or differences in the features that are most influential in determining the two target variables.

The Variables in the dataset are explained as below:

Species: A categorical variable indicating the species of each penguin (Adelie, Gentoo, or Chinstrap)

Island: A categorical variable indicating the island where each penguin was observed (Biscoe, Dream, or Torgersen)

culmen_length_mm: The length of the penguin's culmen (bill) in millimeters.

culmen_depth_mm: The depth of the penguin's culmen (bill) in millimeters.

flipper_length_mm: The length of the penguin's flipper in millimeters

body_mass_g: The mass of the penguin's body in grams.

sex: The sex of the penguin, either male or female

0.3 Summary of work done

The Analysis was done in two parts.

Part I: This part is focused on predicting the Species of the Penguins with the remaining variables as Predictors. We started with Supervised Machine Learning models using Decision Tree, Random Forest, Gaussian Naive Bayes and Support Vector Machines to predict the Species, the important features influencing the prediction, and compared the accuracies obtained from these models.

We then used Unsupervised Machine Learning using Clustering with the K-means Clustering algorithm, which uses the Euclidean distance to cluster the observations. The count of observations obtained from clustering for each of the 3 clusters was compared to the actual counts of the clusters to test for accuracy. The last model used was KNN algorithm with different values of K.

Part II: In this part the same structure of analysis is followed as done in part I, to predict the Sex of the Penguins.

After the data is explored, cleaned, and appropriately encoded where needed, different machine learning models are built. These models are first trained, accuracy and confusion matrix are also analyzed. The results, accuracy and confusion matrix of the model is commented. The features and their importance are also displayed.

Lastly, final conclusions and the Highlights of the study are presented.

0.4 Part I: Predicting the Species of Penguins.

0.4.1 Pre-Processing

```
[65]: import pandas as pd
```

```
[66]: pwd()
```

```
[66]: '/Users/pratik'
```

```
[67]: input_file = ("/Users/pratik/Desktop/Harrisburg University programs/Courses/  
↳Late Fall Courses 2022/ANLY 530 Principles of Machine Learning/Project/  
↳penguins_size.csv")
```

```
[68]: data = pd.read_csv(input_file)
```

```
[69]: data.head()
```

```
[69]:  species      island  culmen_length_mm  culmen_depth_mm  flipper_length_mm  \  
0  Adelie  Torgersen           39.1           18.7           181.0  
1  Adelie  Torgersen           39.5           17.4           186.0  
2  Adelie  Torgersen           40.3           18.0           195.0  
3  Adelie  Torgersen            NaN            NaN            NaN  
4  Adelie  Torgersen           36.7           19.3           193.0  
  
   body_mass_g  sex  
0       3750.0  MALE  
1       3800.0  FEMALE  
2       3250.0  FEMALE  
3          NaN    NaN  
4       3450.0  FEMALE
```

```
[70]: data.dtypes
```

```
[70]: species      object  
island      object  
culmen_length_mm  float64  
culmen_depth_mm  float64  
flipper_length_mm  float64  
body_mass_g      float64  
sex              object  
dtype: object
```

```
[71]: data.describe()
```

```
[71]:      culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  
count      342.000000      342.000000      342.000000      342.000000  
mean        43.921930        17.151170        200.915205      4201.754386
```

std	5.459584	1.974793	14.061714	801.954536
min	32.100000	13.100000	172.000000	2700.000000
25%	39.225000	15.600000	190.000000	3550.000000
50%	44.450000	17.300000	197.000000	4050.000000
75%	48.500000	18.700000	213.000000	4750.000000
max	59.600000	21.500000	231.000000	6300.000000

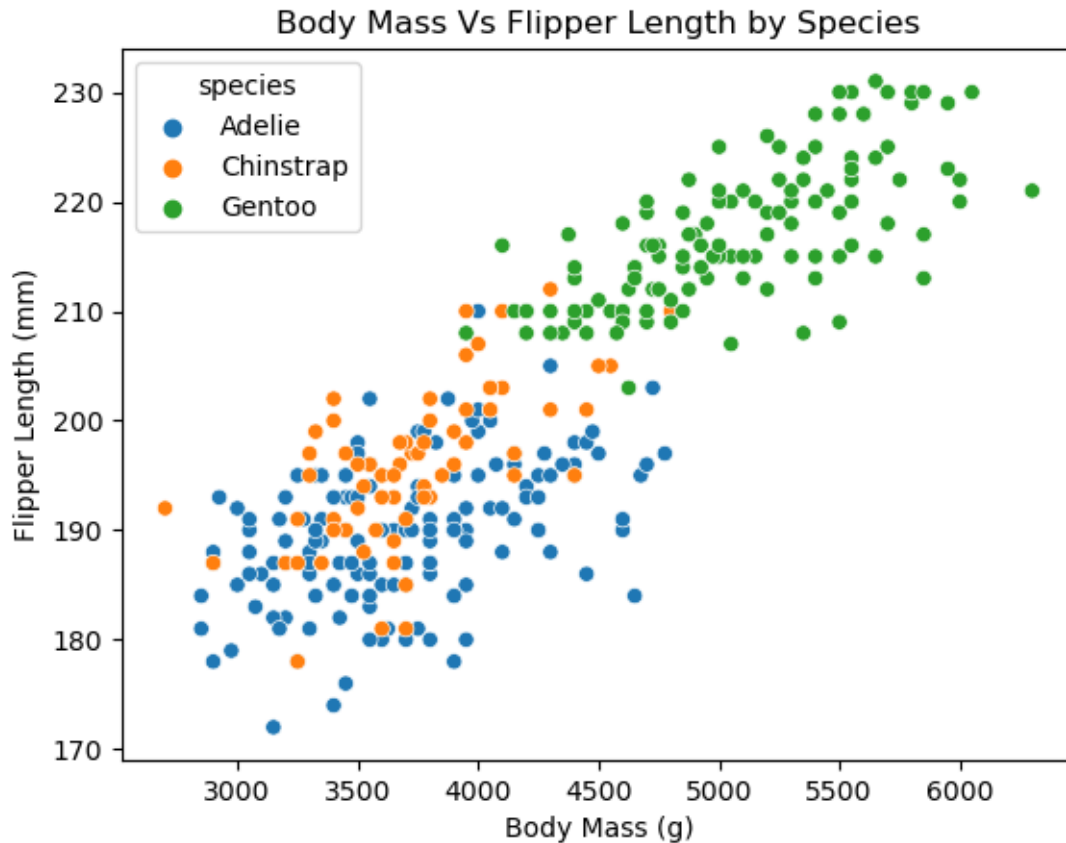
```
[72]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 344 entries, 0 to 343
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   species               344 non-null   object
1   island                344 non-null   object
2   culmen_length_mm      342 non-null   float64
3   culmen_depth_mm       342 non-null   float64
4   flipper_length_mm     342 non-null   float64
5   body_mass_g           342 non-null   float64
6   sex                   334 non-null   object
dtypes: float64(4), object(3)
memory usage: 18.9+ KB
```

0.5 Exploratory Data Analysis

```
[73]: import seaborn as sns
plot = sns.scatterplot(x = "body_mass_g", y = "flipper_length_mm", data = data,
    ↪ hue = "species")
plot.set(xlabel = "Body Mass (g)", ylabel = "Flipper Length (mm)", title = 'Body_
    ↪ Mass Vs Flipper Length by Species')
```


```
[73]: [Text(0.5, 0, 'Body Mass (g)'),
Text(0, 0.5, 'Flipper Length (mm)'),
Text(0.5, 1.0, 'Body Mass Vs Flipper Length by Species')]
```



It can be seen that the Gentoo species are heavier than the other species with a longer Flipper Length as well.

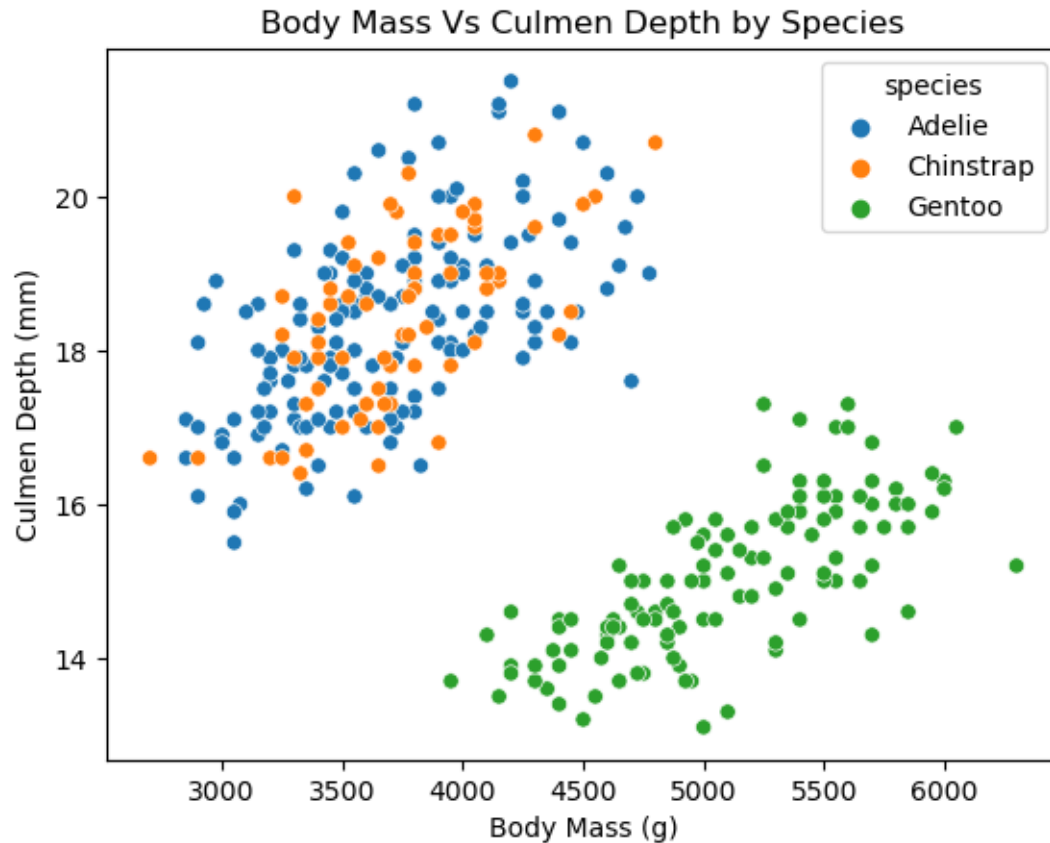
```
[74]: import seaborn as sns
plot = sns.scatterplot(x = "body_mass_g", y = "flipper_length_mm", data = data,
    hue = "sex")
plot.set(xlabel = "Body Mass (g)", ylabel = "Flipper Length (mm)", title = 'Body
    Mass Vs Flipper Length by Sex')
```

```
[74]: [Text(0.5, 0, 'Body Mass (g)'),
Text(0, 0.5, 'Flipper Length (mm)'),
Text(0.5, 1.0, 'Body Mass Vs Flipper Length by Sex')]
```



```
[75]: import seaborn as sns
plot = sns.scatterplot(x = "body_mass_g", y = "culmen_depth_mm", data = data,
    ↪ hue = "species")
plot.set(xlabel = "Body Mass (g)", ylabel = "Culmen Depth (mm)", title = 'Body_
    ↪ Mass Vs Culmen Depth by Species')
```


```
[75]: [Text(0.5, 0, 'Body Mass (g)'),
Text(0, 0.5, 'Culmen Depth (mm)'),
Text(0.5, 1.0, 'Body Mass Vs Culmen Depth by Species')]
```



Its interesting to see that Gentoo has higher body mass compared to other species but lower Culmen Depth


```
[76]: import seaborn as sns
plot = sns.scatterplot(x = "body_mass_g", y = "culmen_depth_mm", data = data,
    ↪ hue = "sex")
plot.set(xlabel = "Body Mass (g)", ylabel = "Culmen Depth (mm)", title = 'Body_
    ↪ Mass Vs Culmen Depth by Sex')
```

```
[76]: [Text(0.5, 0, 'Body Mass (g)'),
Text(0, 0.5, 'Culmen Depth (mm)'),
Text(0.5, 1.0, 'Body Mass Vs Culmen Depth by Sex')]
```



```
[77]: import seaborn as sns
plot = sns.scatterplot(x = "culmen_length_mm", y = "culmen_depth_mm", data = data,
    hue = "species")
plot.set(xlabel = "Culmen Length (m)", ylabel = "Culmen Depth (mm)", title =
    'Culmen Length Vs Culmen Depth by Species')
```


```
[77]: [Text(0.5, 0, 'Culmen Length (m)'),
Text(0, 0.5, 'Culmen Depth (mm)'),
Text(0.5, 1.0, 'Culmen Length Vs Culmen Depth by Species')]
```

The Culmen lengths for Adelie appear to be smaller than the other two species and the Culmen depths for Gentoo appears to be smaller compared to that for Adelie and Chinstrap species

```
[78]: import seaborn as sns
plot = sns.scatterplot(x = "culmen_length_mm", y = "culmen_depth_mm", data = data, hue = "sex")
plot.set(xlabel = "Culmen Length (mm)", ylabel = "Culmen Depth (mm)", title = "Culmen Length Vs Culmen Depth by Sex")
```

```
[78]: [Text(0.5, 0, 'Culmen Length (mm)'),
Text(0, 0.5, 'Culmen Depth (mm)'),
Text(0.5, 1.0, 'Culmen Length Vs Culmen Depth by Sex')]
```



```
[ ]:
```

Culmen Depth and Body Mass for Females appear to be less than that for males

Dropping na's

```
[79]: data.isna().sum()
```

```
[79]: species      0
      island      0
      culmen_length_mm  2
      culmen_depth_mm  2
      flipper_length_mm  2
      body_mass_g    2
      sex         10
      dtype: int64
```

```
[80]: data = data.dropna()
```

```
[81]: data.isna().sum()
```

```
[81]: species          0
      island          0
      culmen_length_mm 0
      culmen_depth_mm  0
      flipper_length_mm 0
      body_mass_g       0
      sex              0
      dtype: int64
```

Checking if any of the categorical data columns has any other values or special characters in columns other than the categories

```
[82]: data.sex.unique()
```

```
[82]: array(['MALE', 'FEMALE', '.'], dtype=object)
```

Columns sex has “.” as an entry, Removing the row with “.” from the sex column.

```
[83]: data.drop(data[(data['sex'] == ".")].index, inplace = True)
```

```
[84]: data.sex.unique()
```

```
[84]: array(['MALE', 'FEMALE'], dtype=object)
```

```
[85]: data.species.unique()
```

```
[85]: array(['Adelie', 'Chinstrap', 'Gentoo'], dtype=object)
```

```
[86]: data.island.unique()
```

```
[86]: array(['Torgersen', 'Biscoe', 'Dream'], dtype=object)
```

```
[87]: data = data.dropna()
```

```
[88]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 333 entries, 0 to 343
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   species              333 non-null   object
1   island               333 non-null   object
2   culmen_length_mm     333 non-null   float64
3   culmen_depth_mm      333 non-null   float64
4   flipper_length_mm    333 non-null   float64
5   body_mass_g          333 non-null   float64
6   sex                  333 non-null   object
```

```
dtypes: float64(4), object(3)
memory usage: 20.8+ KB
```

The observations were 344 before removing na's, 334 after removing na's and 333 after removing “.”

```
[89]: data.describe()
```

```
[89]:
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g
count	333.000000	333.000000	333.000000	333.000000
mean	43.992793	17.164865	200.966967	4207.057057
std	5.468668	1.969235	14.015765	805.215802
min	32.100000	13.100000	172.000000	2700.000000
25%	39.500000	15.600000	190.000000	3550.000000
50%	44.500000	17.300000	197.000000	4050.000000
75%	48.600000	18.700000	213.000000	4775.000000
max	59.600000	21.500000	231.000000	6300.000000

One Hot Encoding all categorical variables (using `pd.get_dummies` method) except species, which is the target variable.

```
[90]: island = pd.get_dummies(data["island"], drop_first = True)
island.head()
```

```
[90]:
```

	Dream	Torgersen
0	0	1
1	0	1
2	0	1
4	0	1
5	0	1

```
[91]: sex = pd.get_dummies(data["sex"], drop_first = True)
sex.head()
```

```
[91]:
```

	MALE
0	1
1	0
2	0
4	0
5	1

Including the dummy columns in the dataframe and creating a new dataframe.

```
[92]: new_data = pd.concat([data,island,sex], axis = 1)
```

```
[93]: new_data.head()
```

```
[93]: species      island  culmen_length_mm  culmen_depth_mm  flipper_length_mm  \
0  Adelie  Torgersen           39.1           18.7           181.0
1  Adelie  Torgersen           39.5           17.4           186.0
2  Adelie  Torgersen           40.3           18.0           195.0
4  Adelie  Torgersen           36.7           19.3           193.0
5  Adelie  Torgersen           39.3           20.6           190.0

      body_mass_g      sex  Dream  Torgersen  MALE
0         3750.0    MALE      0          1      1
1         3800.0  FEMALE      0          1      0
2         3250.0  FEMALE      0          1      0
4         3450.0  FEMALE      0          1      0
5         3650.0    MALE      0          1      1
```

Now we can remove the island and sex columns as we have the dummies. We did not get dummies for Species as we wanted to retain that column. We will later assign numerical values to it using map function. Removing island and sex columns as below.

```
[94]: new_data.drop(['sex', 'island'], axis = 1, inplace = True)
```

```
[95]: new_data.head()
```

```
[95]: species  culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  \
0  Adelie           39.1           18.7           181.0         3750.0
1  Adelie           39.5           17.4           186.0         3800.0
2  Adelie           40.3           18.0           195.0         3250.0
4  Adelie           36.7           19.3           193.0         3450.0
5  Adelie           39.3           20.6           190.0         3650.0

      Dream  Torgersen  MALE
0         0          1      1
1         0          1      0
2         0          1      0
4         0          1      0
5         0          1      1
```

```
[96]: new_data.shape
```

```
[96]: (333, 8)
```

Assigning species as target variable and converting target variable to numeric

0.5.1 Randomizing

```
[97]: import random
target = new_data["species"]
random.seed(12345)
indx = random.sample(range(0, 333), 333)
new_data_rand = new_data.iloc[indx]
target_rand = target.iloc[indx]
```

```
[98]: target_rand.unique()
```

```
[98]: array(['Chinstrap', 'Adelie', 'Gentoo'], dtype=object)
```

```
[99]: target_rand = target_rand.map({'Adelie' : 0, 'Chinstrap' : 1, 'Gentoo' : 2})
target_rand.unique()
```

```
[99]: array([1, 0, 2])
```

0.6 Decision Tree Classifier

Assinging target to Y, predictors to X and splitting into Train and Test in the ratio of 70:30

```
[100]: Y = target_rand
X = new_data_rand.drop(["species"], axis = 1)
```

```
[101]: X.head()
```

```
[101]:      culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  Dream  \
219                50.2                18.7                198.0        3775.0      1
6                  38.9                17.8                181.0        3625.0      0
158                46.1                18.2                178.0        3250.0      1
194                50.9                19.1                196.0        3550.0      1
105                39.7                18.9                184.0        3550.0      0
```

```
      Torgersen  MALE
219           0     0
6             1     0
158           0     0
194           0     1
105           0     1
```

```
[102]: from sklearn.model_selection import train_test_split
Y = target_rand
X = new_data_rand.drop(["species"], axis = 1)
```

```
[103]: target_rand.value_counts()
```

```
[103]: 0    146
      2    119
      1     68
      Name: species, dtype: int64
```

```
[104]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.30,
      ↪random_state = 52)
```

```
[105]: Y_train.value_counts() / Y_train.shape
```

```
[105]: 0    0.442060
      2    0.334764
      1    0.223176
      Name: species, dtype: float64
```

```
[106]: Y_test.value_counts() / Y_test.shape
```

```
[106]: 0    0.43
      2    0.41
      1    0.16
      Name: species, dtype: float64
```

```
[107]: from sklearn import tree
      from sklearn.tree import DecisionTreeClassifier
```

```
[38]: model = tree.DecisionTreeClassifier()
      model = model.fit(X_train, Y_train)
```

Model Evaluation

```
[39]: from sklearn.metrics import confusion_matrix
      from sklearn.metrics import accuracy_score
```

```
[40]: Y_predict = model.predict(X_test)
```

```
[41]: print(confusion_matrix(Y_test, Y_predict))
```

```
[[42  0  1]
 [ 0 16  0]
 [ 0  1 40]]
```

```
[42]: print(accuracy_score(Y_predict, Y_test)*100)
```

98.0

0.7 Random Forest Classifier

```
[108]: from sklearn.ensemble import RandomForestClassifier
print("X_train shape is : ", X_train.shape)
print("X_test shape is : ", X_test.shape)
print("Y_train shape is : ", Y_train.shape)
print("Y_test shape is : ", Y_test.shape)
```

```
X_train shape is : (233, 7)
X_test shape is : (100, 7)
Y_train shape is : (233,)
Y_test shape is : (100,)
```

```
[109]: clf = RandomForestClassifier()
model = clf.fit(X_train, Y_train)
```

Model Evaluation

```
[110]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
Y_predict = clf.predict(Y_test)
```

```
/Users/pratik/opt/anaconda3/lib/python3.9/site-packages/sklearn/base.py:450:
UserWarning: X does not have valid feature names, but RandomForestClassifier was
fitted with feature names
  warnings.warn(
```

```
-----
ValueError                                Traceback (most recent call last)
/var/folders/fp/szpx3gxd4ln849f9dbx_lqw00000gn/T/ipykernel_954/1026824110.py in
↳<module>
      1 from sklearn.metrics import confusion_matrix
      2 from sklearn.metrics import accuracy_score
----> 3 Y_predict = clf.predict(Y_test)

~/opt/anaconda3/lib/python3.9/site-packages/sklearn/ensemble/_forest.py in
↳predict(self, X)
    806         The predicted classes.
    807         """
--> 808         proba = self.predict_proba(X)
    809
    810         if self.n_outputs_ == 1:

~/opt/anaconda3/lib/python3.9/site-packages/sklearn/ensemble/_forest.py in
↳predict_proba(self, X)
    848         check_is_fitted(self)
    849         # Check data
--> 850         X = self._validate_X_predict(X)
```



```

851
852         # Assign chunk of trees to jobs

~/opt/anaconda3/lib/python3.9/site-packages/sklearn/ensemble/_forest.py in
-> _validate_X_predict(self, X)
    577         Validate X whenever one tries to predict, apply, predict_proba.
-> """
    578         check_is_fitted(self)
-> 579         X = self._validate_data(X, dtype=DTYPE, accept_sparse="csr",
-> reset=False)
    580         if issparse(X) and (X.indices.dtype != np.intc or X.indptr.dtype
-> != np.intc):
    581             raise ValueError("No support for np.int64 index based sparse
-> matrices")

~/opt/anaconda3/lib/python3.9/site-packages/sklearn/base.py in
-> _validate_data(self, X, y, reset, validate_separately, **check_params)
    564         raise ValueError("Validation should be done on X, y or both
-> ")
    565         elif not no_val_X and no_val_y:
-> 566             X = check_array(X, **check_params)
    567             out = X
    568         elif no_val_X and not no_val_y:

~/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/validation.py in
-> check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy,
-> force_all_finite, ensure_2d, allow_nd, ensure_min_samples,
-> ensure_min_features, estimator)
    767         # If input is 1D raise error
    768         if array.ndim == 1:
-> 769             raise ValueError(
    770                 "Expected 2D array, got 1D array instead:\narray={}
-> \n"
    771                 "Reshape your data either using array.reshape(-1, 1)
-> if "

ValueError: Expected 2D array, got 1D array instead:
array=[0. 2. 1. 2. 2. 1. 0. 1. 0. 0. 1. 2. 0. 2. 2. 1. 2. 0. 2. 2. 2. 1. 0. 2.
0. 1. 0. 0. 1. 0. 0. 2. 2. 2. 1. 2. 2. 2. 2. 2. 0. 0. 1. 0. 2. 0. 0. 2.
0. 0. 0. 1. 2. 0. 0. 0. 0. 2. 2. 2. 2. 2. 2. 1. 2. 0. 2. 0. 2. 2. 0. 2. 2.
1. 0. 2. 0. 0. 2. 0. 1. 0. 2. 2. 0. 0. 0. 0. 0. 2. 1. 0. 0. 0. 2. 2. 1.
2. 0. 0. 0.].
Reshape your data either using array.reshape(-1, 1) if your data has a single
-> feature or array.reshape(1, -1) if it contains a single sample.

```

```
[46]: print(confusion_matrix(Y_test, Y_predict))
```

```
[[42  0  1]
 [ 0 16  0]
 [ 0  1 40]]
```

```
[47]: print(accuracy_score(Y_test, Y_predict)*100)
```

```
98.0
```

0.7.1 Importance of Features

```
[111]: import pandas as pd
feature_importances = pd.DataFrame(clf.feature_importances_, index = X.columns,
    ↪columns=['importance']).sort_values('importance', ascending=False)
feature_importances
```

```
[111]:
```

	importance
culmen_length_mm	0.361656
flipper_length_mm	0.191926
culmen_depth_mm	0.178203
body_mass_g	0.126603
Dream	0.108121
Torgersen	0.027580
MALE	0.005911

The above results show there was no change in accuracy when we used Random Forest algorithm after first using the Decision Tree Classifier. Both showed an accuracy of 98%. Culmen length, Flipper Length, and Culmen Depth were the top three most important features in predicting the Species of Penguins, in that order.

0.8 Gaussian Naive Bayes Classifier

```
[49]: from sklearn.model_selection import train_test_split

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.30,
    ↪random_state = 52)
Y_train.value_counts() / Y_train.shape
```

```
[49]: 0    0.442060
      2    0.334764
      1    0.223176
      Name: species, dtype: float64
```

```
[50]: Y_test.value_counts() / Y_test.shape
```

```
[50]: 0    0.43
      2    0.41
      1    0.16
```

Name: species, dtype: float64

```
[51]: from sklearn.naive_bayes import GaussianNB
      gnb = GaussianNB()
      model = gnb.fit(X_train, Y_train)
```

Model Evaluation

```
[52]: from sklearn.metrics import confusion_matrix
      from sklearn.metrics import accuracy_score
      Y_predict = model.predict(X_test)
      print(confusion_matrix(Y_test, Y_predict))
```

```
[[33 10  0]
 [ 0 16  0]
 [ 0  0 41]]
```

```
[53]: print(accuracy_score(Y_test, Y_predict)*100)
```

90.0

Above results show that Gaussian Naive Bayes gives us an accuracy of 90% which is less than the 98% accuracy obtained for both the Decision Tree and Random Forest Classifiers.

0.9 Support Vector Machines

```
[54]: new_data_rand.head()
```

```
[54]:      species  culmen_length_mm  culmen_depth_mm  flipper_length_mm  \
219  Chinstrap             50.2             18.7             198.0
  6    Adelie              38.9             17.8             181.0
158  Chinstrap             46.1             18.2             178.0
194  Chinstrap             50.9             19.1             196.0
105  Adelie               39.7             18.9             184.0

      body_mass_g  Dream  Torgersen  MALE
219      3775.0      1         0       0
  6      3625.0      0         1       0
158      3250.0      1         0       0
194      3550.0      1         0       1
105      3550.0      0         0       1
```

```
[55]: Y = target_rand
      X = new_data_rand.drop(["species"], axis = 1)
```

```
[56]: Y.head()
```

```
[56]: 219    1
      6      0
      158    1
      194    1
      105    0
      Name: species, dtype: int64
```

Normalizing the X variables

```
[57]: from sklearn.preprocessing import MinMaxScaler
      scaler = MinMaxScaler()
      scaled = scaler.fit_transform(X)
      print(scaled)
```

```
[[0.65818182 0.66666667 0.44067797 ... 1.         0.         0.         ]
 [0.24727273 0.55952381 0.15254237 ... 0.         1.         0.         ]
 [0.50909091 0.60714286 0.10169492 ... 1.         0.         0.         ]
 ...
 [0.41454545 0.25         0.69491525 ... 0.         0.         0.         ]
 [0.48363636 0.66666667 0.27118644 ... 1.         0.         0.         ]
 [0.16         0.58333333 0.16949153 ... 1.         0.         0.         ]]
```

Update new scaled X for data modelling

```
[58]: from sklearn.model_selection import train_test_split
      Y = target_rand
      X = scaled
      X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size = 0.3,
      ↪random_state = 52)
```

```
[59]: Y_train.value_counts() / Y_train.shape
```

```
[59]: 0    0.442060
      2    0.334764
      1    0.223176
      Name: species, dtype: float64
```

```
[60]: Y_test.value_counts() / Y_test.shape
```

```
[60]: 0    0.43
      2    0.41
      1    0.16
      Name: species, dtype: float64
```

```
[61]: X_train.shape
```

```
[61]: (233, 7)
```

```
[62]: Y_train.shape
```

```
[62]: (233,)
```

```
[63]: X_test.shape
```

```
[63]: (100, 7)
```

```
[64]: Y_test.shape
```

```
[64]: (100,)
```

0.9.1 Designing the model

```
[65]: from sklearn import svm  
      clf = svm.SVC(kernel = "linear")  
      model = clf.fit(X_train, Y_train)
```

```
[66]: Y_predict = model.predict(X_test)
```

Model Evaluation

```
[67]: from sklearn.metrics import confusion_matrix  
      from sklearn.metrics import accuracy_score  
      print(confusion_matrix(Y_test, Y_predict))
```

```
[[43  0  0]  
 [ 0 16  0]  
 [ 0  0 41]]
```

```
[68]: print(accuracy_score(Y_test, Y_predict)*100)
```

```
100.0
```

We get an accuracy of 100% using the Linear Kernel. This model correctly classified all the Species of Penguins.

0.9.2 Using RBF Kernel

```
[69]: clf = svm.SVC(kernel='rbf', degree = 8)  
      model = clf.fit(X_train, Y_train)  
      Y_predict = model.predict(X_test)  
      print(confusion_matrix(Y_test, Y_predict))
```

```
[[43  0  0]  
 [ 0 16  0]  
 [ 0  0 41]]
```

```
[70]: print(accuracy_score(Y_test, Y_predict)*100)
```

100.0

0.9.3 Using polynomial Kernel

```
[71]: clf = svm.SVC(kernel='poly', gamma =0.3)
      model = clf.fit(X_train, Y_train)
      Y_predict = model.predict(X_test)
      print(confusion_matrix(Y_test, Y_predict))
```

```
[[43  0  0]
 [ 1 15  0]
 [ 0  0 41]]
```

```
[72]: print(accuracy_score(Y_test, Y_predict)*100)
```

99.0

As seen above, we got an accuracy of 100% with the Linear kernel and RBF kernel(Non-Linear) for Support Vector Machine classifiers and 99% for the Polynomial kernel(Non-Linear).

We used Decision Tree, Random Forest, Gaussian Naive Bayes and Support Vector Machine models for predicting the species of penguins. Out of the three models used, two models gave an accuracy of 98% with the exception of Gaussian Naive Bayes which predicted the species with a lower accuracy of 90%, compared to the other three models. Support Vector Machines gave an accuracy of 100% for the Linear and RBF kernel and 99% when using the polynomial kernel.

0.10 Unsupervised: Clustering using the K-means algorithm

```
[73]: new_data_rand.head()
```

```
[73]:
```

	species	culmen_length_mm	culmen_depth_mm	flipper_length_mm	\
219	Chinstrap	50.2	18.7	198.0	
6	Adelie	38.9	17.8	181.0	
158	Chinstrap	46.1	18.2	178.0	
194	Chinstrap	50.9	19.1	196.0	
105	Adelie	39.7	18.9	184.0	

	body_mass_g	Dream	Torgersen	MALE
219	3775.0	1	0	0
6	3625.0	0	1	0
158	3250.0	1	0	0
194	3550.0	1	0	1
105	3550.0	0	0	1

Normalizing

```
[75]: Y = target_rand
      X = new_data_rand.drop(["species"], axis = 1)
      X.head()
```

```
[75]:
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g	Dream	\
219	50.2	18.7	198.0	3775.0	1	
6	38.9	17.8	181.0	3625.0	0	
158	46.1	18.2	178.0	3250.0	1	
194	50.9	19.1	196.0	3550.0	1	
105	39.7	18.9	184.0	3550.0	0	

	Torgersen	MALE
219	0	0
6	1	0
158	0	0
194	0	1
105	0	1

```
[76]: Y.head()
```

```
[76]: 219    1
      6     0
      158   1
      194   1
      105   0
      Name: species, dtype: int64
```

```
[77]: from sklearn.preprocessing import MinMaxScaler
      scaler = MinMaxScaler()
      scaled = scaler.fit_transform(X)
      print(scaled)
```

```
[[0.65818182 0.66666667 0.44067797 ... 1.         0.         0.         ]
 [0.24727273 0.55952381 0.15254237 ... 0.         1.         0.         ]
 [0.50909091 0.60714286 0.10169492 ... 1.         0.         0.         ]
 ...
 [0.41454545 0.25         0.69491525 ... 0.         0.         0.         ]
 [0.48363636 0.66666667 0.27118644 ... 1.         0.         0.         ]
 [0.16        0.58333333 0.16949153 ... 1.         0.         0.         ]]
```


```
[78]: Y = target_rand
      X = scaled
```

Determining the best value of K clusters using the predictor variables and Kmeans Clustering and random seed of 12345 .(Note: species as Y=target is mapped to numeric, however species in new_data is not.)

```
[79]: from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

sse = {}
last_sse = 17592402.70373319
for k in range(1, 20):
    kmeans = KMeans(n_clusters=k, random_state=12345, n_init = 25).fit(X) ##
    ↪Note X is scaled
    #print(data["clusters"])
    sse[k] = kmeans.inertia_ # Inertia: Sum of distances of samples to their
    ↪closest cluster center
    change_per = (last_sse-kmeans.inertia_)/last_sse*100
    print ('At k= ',k,'The percentage of change in SSE is ',change_per,'%')
    last_sse = kmeans.inertia_
plt.figure()
plt.plot(list(sse.keys()), list(sse.values()))
plt.xlabel("Number of cluster")
plt.ylabel("SSE")
plt.show()
```

```
At k= 1 The percentage of change in SSE is 99.99847719635022 %
At k= 2 The percentage of change in SSE is 34.438660476557246 %
At k= 3 The percentage of change in SSE is 33.20816022869925 %
At k= 4 The percentage of change in SSE is 29.775775145226262 %
At k= 5 The percentage of change in SSE is 40.44501333768008 %
At k= 6 The percentage of change in SSE is 26.435420116534058 %
At k= 7 The percentage of change in SSE is 31.25444566076953 %
At k= 8 The percentage of change in SSE is 37.22018068641265 %
At k= 9 The percentage of change in SSE is 17.967286903698497 %
At k= 10 The percentage of change in SSE is 15.891489060276145 %
At k= 11 The percentage of change in SSE is 7.297052685601946 %
At k= 12 The percentage of change in SSE is 4.281490997278908 %
At k= 13 The percentage of change in SSE is 3.098632216159535 %
At k= 14 The percentage of change in SSE is 3.6327810730025667 %
At k= 15 The percentage of change in SSE is 4.189420625521325 %
At k= 16 The percentage of change in SSE is 2.879518781607917 %
At k= 17 The percentage of change in SSE is 3.383907108419674 %
At k= 18 The percentage of change in SSE is 2.3018622359209813 %
At k= 19 The percentage of change in SSE is 3.3416461490160887 %
```

Two elbows can be seen in the graph from one from $K = 3$ to $K = 5$ and the other from $k = 5$ to $k = 8$. We can say that K in this range (3 to 8) seems like a good fit as the change in percentage of sse is more compared to the values above $K = 8$, within group error is considerably less compared to values below $K = 3$ and the between group distance looks acceptable.

```
[83]: from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=3, random_state=1234, n_init = 25 ).fit(X)
kmeans.cluster_centers_
```

```
[83]: array([[ 4.40798226e-01,  6.23790166e-01,  3.59101557e-01,
                2.83028455e-01,  1.00000000e+00,  1.94289029e-16,
                5.04065041e-01],
               [ 4.94614065e-01,  4.88544474e-01,  6.32235369e-01,
                6.03446017e-01, -1.66533454e-16,  2.16981132e-01,
                1.00000000e+00],
               [ 3.59265734e-01,  3.13759158e-01,  5.02933507e-01,
                3.90625000e-01, -1.66533454e-16,  2.30769231e-01,
                5.55111512e-16]])
```

```
[84]: pd.Series(kmeans.labels_).value_counts()
```

```
[84]: 0    123
      1    106
      2    104
      dtype: int64
```

```
[90]: target_rand.value_counts()
```

```
[90]: 0    146
      2    119
      1     68
      Name: species, dtype: int64
```

```
[92]: print(confusion_matrix(kmeans.labels_, target_rand))
```

```
[[55 68  0]
 [45  0 61]
 [46  0 58]]
```

```
[93]: print(accuracy_score(kmeans.labels_, target_rand)*100)
```

```
33.93393339333936
```

The accuracy score was found to be less when comparing the clusters obtained from K-means algorithm. One of the reasons responsible for some misclassifications could be that the predicted K_means_ clusters has the species with the second highest counts labelled as 1 and the second highest counts in target_rand is labelled as 2.

0.11 Supervised: KNN Classification algorithm

```
[562]: Y = target_rand
      X = scaled
      X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size = 0.3,
      ↪random_state = 52)
```

0.11.1 KNN Classification model with K = 3

```
[563]: from sklearn.neighbors import KNeighborsClassifier
      neigh = KNeighborsClassifier(n_neighbors=3)
      model = neigh.fit(X, Y)
      Y_predict = model.predict(X_test)
```

/Users/pratik/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be

accepted. Set `keepdims` to True or False to avoid this warning.
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)

```
[564]: from sklearn.metrics import confusion_matrix
       from sklearn.metrics import accuracy_score

       print(confusion_matrix(Y_test, Y_predict))
```

```
[[43  0  0]
 [ 0 16  0]
 [ 0  0 41]]
```

```
[565]: print(accuracy_score(Y_test, Y_predict)*100)
```

100.0

We get an accuracy of 100 % with K=3

0.11.2 Model with K=5

```
neigh = KNeighborsClassifier(n_neighbors=5) model = neigh.fit(X, Y) Y_predict =
model.predict(X_test) print(confusion_matrix(Y_test, Y_predict))
```

```
[567]: Y_predict
```

```
[567]: array([0, 2, 1, 2, 2, 1, 0, 1, 0, 0, 1, 2, 0, 2, 2, 1, 2, 0, 2, 2, 2, 1,
            0, 2, 0, 1, 0, 0, 1, 0, 0, 2, 2, 2, 1, 2, 2, 2, 2, 2, 0, 0, 1, 0,
            2, 0, 0, 2, 0, 0, 0, 1, 2, 0, 0, 0, 0, 2, 2, 2, 2, 2, 1, 2, 0, 2,
            0, 2, 2, 0, 2, 2, 1, 0, 2, 0, 0, 2, 0, 1, 0, 2, 2, 0, 0, 0, 0, 0,
            2, 1, 0, 0, 0, 2, 2, 1, 2, 0, 0, 0])
```

```
[568]: print(accuracy_score(Y_test, Y_predict)*100)
```

100.0

0.12 Model with K = 8

```
[569]: neigh = KNeighborsClassifier(n_neighbors=8)
       model = neigh.fit(X, Y)
       Y_predict = model.predict(X_test)
       print(confusion_matrix(Y_test, Y_predict))
```

```
[[43  0  0]
 [ 0 16  0]
 [ 0  0 41]]
```

/Users/pratik/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode`

typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
[570]: print(accuracy_score(Y_test, Y_predict)*100)
```

```
100.0
```

All values of K used gave an accuracy of 100% in predicting the Species of the Penguins.

0.12.1 Model I: Conclusion

For predicting the Species of the penguins, both the KNN algorithm and Support Vector Machines (Linear and RBF) gave the maximum accuracy at 100% followed by the Polynomial Kernel SVM at 99%. The Random Forest and Decision Tree Classifiers both predicted with an accuracy both at 97% and the Gaussian Naive bayes gave us the least accuracy at 90%. Random Forest predicted Culmen length, Flipper Length, and Culmen Depth as the top three important features in predicting the Species, in that order.

0.13 Model II: Predicting the sex of Penguins

```
[129]: data.head()
```

```
[129]:  species      island  culmen_length_mm  culmen_depth_mm  flipper_length_mm  \
0  Adelie  Torgersen         39.1             18.7             181.0
1  Adelie  Torgersen         39.5             17.4             186.0
2  Adelie  Torgersen         40.3             18.0             195.0
4  Adelie  Torgersen         36.7             19.3             193.0
5  Adelie  Torgersen         39.3             20.6             190.0

      body_mass_g      sex
0         3750.0    MALE
1         3800.0  FEMALE
2         3250.0  FEMALE
4         3450.0  FEMALE
5         3650.0    MALE
```

```
[130]: data.isna().sum()
```

```
[130]: species      0
island          0
culmen_length_mm  0
culmen_depth_mm  0
flipper_length_mm 0
body_mass_g      0
```

```
sex                0
dtype: int64
```

One Hot Encoding Species and Island. not encoding Sex as we need to retain the column in dataframe for analysis

```
[131]: island = pd.get_dummies(data["island"], drop_first = True)
species = pd.get_dummies(data["species"], drop_first = True)
island.head()
```

```
[131]:      Dream  Torgersen
0         0           1
1         0           1
2         0           1
4         0           1
5         0           1
```

```
[132]: species.head()
```

```
[132]:      Chinstrap  Gentoo
0           0         0
1           0         0
2           0         0
4           0         0
5           0         0
```

Including the dummy columns in the dataframe and creating a new dataframe

```
[133]: new_data = pd.concat([data, species, island], axis = 1)
```

```
[134]: new_data.head()
```

```
[134]:   species      island  culmen_length_mm  culmen_depth_mm  flipper_length_mm  \
0  Adelie  Torgersen          39.1             18.7             181.0
1  Adelie  Torgersen          39.5             17.4             186.0
2  Adelie  Torgersen          40.3             18.0             195.0
4  Adelie  Torgersen          36.7             19.3             193.0
5  Adelie  Torgersen          39.3             20.6             190.0

      body_mass_g      sex  Chinstrap  Gentoo  Dream  Torgersen
0       3750.0    MALE           0         0       0           1
1       3800.0  FEMALE           0         0       0           1
2       3250.0  FEMALE           0         0       0           1
4       3450.0  FEMALE           0         0       0           1
5       3650.0    MALE           0         0       0           1
```

Dropping species and island from the dataframe

```
[135]: new_data.drop(['species', 'island'], axis = 1, inplace = True)
new_data.head()
```

```
[135]:      culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  sex \
0           39.1           18.7           181.0          3750.0  MALE
1           39.5           17.4           186.0          3800.0  FEMALE
2           40.3           18.0           195.0          3250.0  FEMALE
4           36.7           19.3           193.0          3450.0  FEMALE
5           39.3           20.6           190.0          3650.0  MALE

      Chinstrap  Gentoo  Dream  Torgersen
0             0       0      0           1
1             0       0      0           1
2             0       0      0           1
4             0       0      0           1
5             0       0      0           1
```

Assigning sex to target and converting target to numeric and randomizing the data.

```
[163]: import random
target = new_data["sex"]
random.seed(12345)
indx = random.sample(range(0, 333), 333)
new_data_rand = new_data.iloc[indx]
target_rand = target.iloc[indx]
```

```
[164]: target.unique()
```

```
[164]: array(['MALE', 'FEMALE'], dtype=object)
```

```
[165]: target_rand = target_rand.map({'MALE' : 0, 'FEMALE' : 1})
target_rand.unique()
```

```
[165]: array([1, 0])
```

0.14 Desision Tree Classifier

```
[166]: Y = target_rand
X = new_data_rand.drop(["sex"],axis = 1)
```

```
[167]: X.head()
```

```
[167]:      culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  \
219           50.2           18.7           198.0          3775.0
6            38.9           17.8           181.0          3625.0
158           46.1           18.2           178.0          3250.0
194           50.9           19.1           196.0          3550.0
```

105	39.7	18.9	184.0	3550.0
	Chinstrap	Gentoo	Dream	Torgersen
219	1	0	1	0
6	0	0	0	1
158	1	0	1	0
194	1	0	1	0
105	0	0	0	0

```
[168]: Y.head()
```

```
[168]: 219    1
        6    1
        158  1
        194  0
        105  0
        Name: sex, dtype: int64
```

```
[169]: from sklearn.model_selection import train_test_split
        X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.30,
        ↪random_state = 52)
```

```
[170]: Y_train.value_counts() / Y_train.shape
```

```
[170]: 0    0.519313
        1    0.480687
        Name: sex, dtype: float64
```

```
[171]: Y_test.value_counts() / Y_test.shape
```

```
[171]: 1    0.53
        0    0.47
        Name: sex, dtype: float64
```

```
[172]: from sklearn import tree
        from sklearn.tree import DecisionTreeClassifier
        tree = DecisionTreeClassifier()
        model = tree.fit(X_train, Y_train)
```

0.14.1 Model Evaluation

```
[173]: from sklearn.metrics import confusion_matrix
        from sklearn.metrics import accuracy_score
        Y_predict = model.predict(X_test)
        print(confusion_matrix(Y_test, Y_predict))
```

```
[[41  6]
 [ 4 49]]
```

```
[174]: print(accuracy_score(Y_test, Y_predict)*100)
```

```
90.0
```

Using Decision Tree gives us an accuracy of 90%. The percentage of False Positives and False Negative classified were fairly close to each other at 4% and 6% respectively.

0.15 Random Forest

```
[178]: from sklearn.ensemble import RandomForestClassifier
      clf = RandomForestClassifier()
      model = clf.fit(X_train, Y_train)
      Y_predict = model.predict(X_test)
      print(confusion_matrix(Y_test, Y_predict))
```

```
[[43  4]
 [ 5 48]]
```

```
[179]: print(accuracy_score(Y_test, Y_predict)*100)
```

```
91.0
```

0.15.1 Importance of Features

```
[177]: import pandas as pd
      feature_importances = pd.DataFrame(clf.feature_importances_, index = X.columns,
      ↪ columns=['importance']).sort_values('importance', ascending=False)
      feature_importances
```

```
[177]:
```

	importance
culmen_depth_mm	0.292731
body_mass_g	0.284620
culmen_length_mm	0.227383
flipper_length_mm	0.131163
Gentoo	0.027388
Chinstrap	0.016665
Dream	0.011612
Torgersen	0.008439

The most important features in predicting the Sex of the Penguins were the Culmen_depth, body_mass and Culmen_length in that order.

Using Random Forest increased the accuracy to 91%

0.16 Gaussian Naive Bayes Algorithm

```
[180]: from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
model = gnb.fit(X_train, Y_train)
Y_predict = model.predict(X_test)
print(confusion_matrix(Y_test, Y_predict))
```

```
[[30 17]
 [20 33]]
```

```
[181]: print(accuracy_score(Y_test, Y_predict)*100)
```

63.0

The accuracy decreased to 63 % by using the Gaussian Naive Bayes classifier

0.17 Support Vector Machines

Normalizing the X variables

```
[182]: X.head()
```

```
[182]:      culmen_length_mm  culmen_depth_mm  flipper_length_mm  body_mass_g  \
219                50.2                18.7                198.0        3775.0
6                  38.9                17.8                181.0        3625.0
158                46.1                18.2                178.0        3250.0
194                50.9                19.1                196.0        3550.0
105                39.7                18.9                184.0        3550.0
```

```
      Chinstrap  Gentoo  Dream  Torgersen
219           1        0        1          0
6             0        0        0          1
158           1        0        1          0
194           1        0        1          0
105           0        0        0          0
```

```
[183]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
scaled = scaler.fit_transform(X)
print(scaled)
```

```
[[0.65818182 0.66666667 0.44067797 ... 0.         1.         0.         ]
 [0.24727273 0.55952381 0.15254237 ... 0.         0.         1.         ]
 [0.50909091 0.60714286 0.10169492 ... 0.         1.         0.         ]
 ...
 [0.41454545 0.25         0.69491525 ... 1.         0.         0.         ]
 [0.48363636 0.66666667 0.27118644 ... 0.         1.         0.         ]
 [0.16        0.58333333 0.16949153 ... 0.         1.         0.         ]]
```

```
[184]: Y = target_rand
X = scaled
X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size = 0.3,
↳random_state = 52)
```

```
[185]: from sklearn import svm
clf = svm.SVC(kernel = "linear")
model = clf.fit(X_train, Y_train)
Y_predict = model.predict(X_test)
print(confusion_matrix(Y_test, Y_predict))
```

```
[[42  5]
 [ 5 48]]
```

```
[186]: print(accuracy_score(Y_test, Y_predict)*100)
```

90.0

As we see the Linear Kernel gives an accuracy of 90%

0.17.1 RBF Kernel

```
[187]: clf = svm.SVC(kernel='rbf', degree = 8)
model = clf.fit(X_train, Y_train)
Y_predict = model.predict(X_test)
print(confusion_matrix(Y_test, Y_predict))
```

```
[[44  3]
 [ 5 48]]
```

```
[188]: print(accuracy_score(Y_test, Y_predict)*100)
```

92.0

Using the RBF Kernel the accuracy increased to 92%

0.17.2 Polynomial Kernel

```
[189]: clf = svm.SVC(kernel='poly', gamma =0.3)
model = clf.fit(X_train, Y_train)
Y_predict = model.predict(X_test)
print(confusion_matrix(Y_test, Y_predict))
```

```
[[31 16]
 [ 3 50]]
```

```
[190]: print(accuracy_score(Y_test, Y_predict)*100)
```

81.0

The polynomail kernel gave the least accuracy between the 3 Kernels used at 81%

0.18 Unsupervised: Clustering using the K-means algorithm

```
[191]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
scaled = scaler.fit_transform(X)
Y = target_rand
X = scaled
print(X)
```


```
[[0.65818182 0.66666667 0.44067797 ... 0.          1.          0.          ]
 [0.24727273 0.55952381 0.15254237 ... 0.          0.          1.          ]
 [0.50909091 0.60714286 0.10169492 ... 0.          1.          0.          ]
 ...
 [0.41454545 0.25          0.69491525 ... 1.          0.          0.          ]
 [0.48363636 0.66666667 0.27118644 ... 0.          1.          0.          ]
 [0.16          0.58333333 0.16949153 ... 0.          1.          0.          ]]
```

```
[192]: from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

sse = {}
last_sse = 17592402.70373319
for k in range(1, 20):
    kmeans = KMeans(n_clusters=k, random_state=12345, n_init = 25).fit(X) ##
    ↪Note X is scaled
    #print(data["clusters"])
    sse[k] = kmeans.inertia_ # Inertia: Sum of distances of samples to their
    ↪closest cluster center
    change_per = (last_sse-kmeans.inertia_)/last_sse*100
    print('At k= ',k,'The percentage of change in SSE is ',change_per,'%')
    last_sse = kmeans.inertia_
plt.figure()
plt.plot(list(sse.keys()), list(sse.values()))
plt.xlabel("Number of cluster")
plt.ylabel("SSE")
plt.show()
```

```
At k= 1 The percentage of change in SSE is 99.998208072535 %
At k= 2 The percentage of change in SSE is 48.81790114012003 %
At k= 3 The percentage of change in SSE is 52.222510950290946 %
At k= 4 The percentage of change in SSE is 45.4073175382075 %
At k= 5 The percentage of change in SSE is 54.12067371485251 %
At k= 6 The percentage of change in SSE is 19.672990727569847 %
At k= 7 The percentage of change in SSE is 12.240089585935465 %
```

At k= 8 The percentage of change in SSE is 9.57089372768644 %
 At k= 9 The percentage of change in SSE is 9.92506387510864 %
 At k= 10 The percentage of change in SSE is 10.269031034868561 %
 At k= 11 The percentage of change in SSE is 8.656628110688978 %
 At k= 12 The percentage of change in SSE is 5.469863889931658 %
 At k= 13 The percentage of change in SSE is 3.9109261926731316 %
 At k= 14 The percentage of change in SSE is 4.289380163025055 %
 At k= 15 The percentage of change in SSE is 3.2092043953886145 %
 At k= 16 The percentage of change in SSE is 2.8537143628810124 %
 At k= 17 The percentage of change in SSE is 1.8882221598651179 %
 At k= 18 The percentage of change in SSE is 4.218518270236141 %
 At k= 19 The percentage of change in SSE is 2.8627988764956878 %



An elbow can be seen in the graph from $K = 3$ to $K = 5$, we can say that K in this range seems like a good fit as the change in percentage of sse is more compared to the values above $K = 5$, within group error is considerably less compared to values below $K = 3$ and the between group distance looks acceptable.

```
[193]: from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=2, random_state=1234, n_init = 25 ).fit(X)
kmeans.cluster_centers_
```

```
[193]: array([[ 5.62475172e-01,  2.25790316e-01,  7.66699900e-01,
              6.64565826e-01,  2.77555756e-17,  1.00000000e+00,
              -1.11022302e-16,  2.22044605e-16],
              [ 3.60169924e-01,  6.27447708e-01,  3.37636623e-01,
              2.81866563e-01,  3.17757009e-01, -3.33066907e-16,
              5.74766355e-01,  2.19626168e-01]])
```

```
[194]: pd.Series(kmeans.labels_).value_counts()
```

```
[194]: 1    214
       0    119
       dtype: int64
```

```
[195]: target_rand.value_counts()
```

```
[195]: 0    168
       1    165
       Name: sex, dtype: int64
```

```
[196]: print(confusion_matrix(kmeans.labels_, target_rand))
```

```
[[ 61  58]
 [107 107]]
```

```
[197]: print(accuracy_score(kmeans.labels_, target_rand)*100)
```

```
50.45045045045045
```

We used $K = 2$ for clustering to test the accuracy of predicting the classes., which came out to be 50.45%

0.18.1 Supervised: Using KNN algorithm

0.18.2 Model with $K = 3$

```
[198]: from sklearn.neighbors import KNeighborsClassifier
       neigh = KNeighborsClassifier(n_neighbors=3)
       model = neigh.fit(X, Y)
       Y_predict = model.predict(X_test)
       print(confusion_matrix(Y_test, Y_predict))
```

```
[[45  2]
 [ 4 49]]
```

/Users/pratik/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which

the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
[199]: print(accuracy_score(Y_test, Y_predict)*100)
```

94.0

Using KNN algorithm gives an accuracy of 94%

0.18.3 Model with K= 5

```
[200]: from sklearn.neighbors import KNeighborsClassifier
neigh = KNeighborsClassifier(n_neighbors=5)
model = neigh.fit(X, Y)
Y_predict = model.predict(X_test)
print(confusion_matrix(Y_test, Y_predict))
```

```
[[45  2]
 [ 3 50]]
```

/Users/pratik/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
[201]: print(accuracy_score(Y_test, Y_predict)*100)
```

95.0

0.18.4 MODEL WITH K = 6

```
[202]: from sklearn.neighbors import KNeighborsClassifier
neigh = KNeighborsClassifier(n_neighbors=6)
model = neigh.fit(X, Y)
Y_predict = model.predict(X_test)
print(confusion_matrix(Y_test, Y_predict))
```

```
[[46  1]
 [ 3 50]]
```

/Users/pratik/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will

change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
[203]: print(accuracy_score(Y_test, Y_predict)*100)
```

96.0

0.18.5 MODEL WITH K = 10

```
[204]: from sklearn.neighbors import KNeighborsClassifier
neigh = KNeighborsClassifier(n_neighbors=10)
model = neigh.fit(X, Y)
Y_predict = model.predict(X_test)
print(confusion_matrix(Y_test, Y_predict))
```

```
[[44  3]
 [ 4 49]]
```

/Users/pratik/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
[205]: print(accuracy_score(Y_test, Y_predict)*100)
```

93.0

0.18.6 MODEL WITH K = 20

```
[206]: from sklearn.neighbors import KNeighborsClassifier
neigh = KNeighborsClassifier(n_neighbors=18)
model = neigh.fit(X, Y)
Y_predict = model.predict(X_test)
print(confusion_matrix(Y_test, Y_predict))
```

```
[[40  7]
 [ 7 46]]
```

/Users/pratik/opt/anaconda3/lib/python3.9/site-packages/sklearn/neighbors/_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which

the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
[207]: print(accuracy_score(Y_test, Y_predict)*100)
```

86.0

The model with $K = 6$ gave us the highest accuracy of 96%. The accuracy increased from $K=3$ to $k=5$ and to $K=6$. However increasing the values of K beyond 7 to 10,18 resulted in decrease in accuracies to 93 % and 86% respectively.

0.19 Part II: Conclusion

In predicting the Sex of Penguins, the accuracy of Decision Tree Classifier was 90% which further increased to 91% by using the Random Forest classifier. The Accuracy decreased to 63% when using the Gaussian Naive Bayes algorithm. Random Forest predicted the Culmen_depth, body_mass and Culmen_length as the top three most important features in that order.

The RBF Kernel in Support vector machines gave the maximum accuracy of 92% amongst all Kernels used, followed closely by Linear Kernel at 90%. Polynomial Kernel resulted in a decrease of accuracy to 81%. We then used Clustering to predict the best value of number of clusters using the Kmeans algorithm. From the Elbow method, value of K in the range of 3 to 5 seemed like a good range. We then used the KNN algorithm to predict the sex which resulted in a maximum accuracy of 96% at $K=6$.

Hence amongst all the algorithms used, KNN algorithm ($k=6$) gave the maximum accuracy at 96%, followed by the Support Vector machines RBF Kernel predicted at 92%, Random Forest at 91% and Decision Tree Classifier at 90%.

0.20 Highlights

It was interesting to see that the Culmen length, Flipper Length, and Culmen Depth were the Top Three most important features in predicting the Species, which changed to Culmen_depth, body_mass and Culmen_length when predicting the Sex of the Penguins.